

Multidataset-Health

Investigating Generalization in Health Risk Prediction
using Multi-Dataset Learning

Members

SSR+

- | | |
|------------------------------|------------|
| 1. Thanagorn Chaiyut | 6631321321 |
| 2. Tinna Nillapong | 6633083321 |
| 3. Sahatsawat Wongpiyabovorn | 6633257221 |

เราทำมากกว่าพืช

- | | |
|--------------------------|------------|
| 1. Saruj Chutapornpong | 6631353421 |
| 2. Theerachot Ounlum | 6631328821 |
| 3. Nattawat Maneesongsri | 6631315621 |

เราทำมากกว่าพืช

Problem Statement

Health prediction models trained on a single dataset **often fail to generalize to other datasets** due to differences in data distribution and feature representation.

What we do is **combine multiple datasets — using both Union and Intersection** feature strategies — to train shared models, then evaluate whether **these models can outperform single-dataset baselines when tested on individual datasets.**

Workflow

1. Data Preparation
2. Training Setup
3. Model Baseline
4. Concatenation Setup
5. Concatenation Result
6. Feature Importance
7. Error Analysis
8. Conclusion

1. Data Preparation

1. Sources

- **Statlog Heart Disease** - Dataset for Heart Disease prediction
- **Coronary Heart Disease** - CHD cases and clinical variables describing patient health conditions.
- **Framingham** - Dataset used to predict 10-year CHD risk.
- **Heart Disease** - Dataset used to predict Heart Disease
- **Stroke** - Dataset used to predict Stroke Risk

2. Dataset Columns

- **Statlog Heart Disease** - age, sex, chest-pain, rest-bp, serum-chol, fasting-blood-sugar, electrocardiographic, max-heart-rate, angina, oldpeak, slope, major-vessels, thal, heart-disease
- **Coronary Heart Disease** - sbp, tobacco, ldl, adiposity, famhist, typea, obesity, alcohol, age, chd

1. Data Preparation

2. Dataset Columns (cont.)

- **Framingham** - Sex, age, education, currentSmoker, cigsPerDay, BPMeds, prevalentStroke, prevalentHyp, diabetes, totChol, sysBP, diaBP, BMI, heartRate, glucose, TenYearCHD
- **Heart Disease** - age, sex, cp, trestbps, chol, fbs, restecg, thalach, exang, oldpeak, slope, ca, thal, target
- **Stroke** - id, gender, age, hypertension, heart_disease, ever_married, work_type, Residence_type, avg_glucose_level, bmi, smoking_status, stroke

3. Data Cleaning

- Rename Columns
- Change Data types
- Map Values (eg. .map() in pandas)
- Handle Missing Values
- Drop Duplicated Rows
- Feature Scaling

2. Training Setup

- **Train Test Split:** 80/20 (stratified)
- **Cross Validation** with StratifiedKFold (k=5) to preserves class distribution in each fold
- **Imbalanced handling** with SMOTE applied when class ratio ≥ 2.0

- **8 Classification Models**

- Logistic Regression
- KNN
- Decision Tree
- SVM
- Naive Bayes
- Random Forest
- XGBoost
- LightGBM


Model	Loss Function
Logistic Regression	Binary Cross-Entropy (Log Loss)
K-Nearest Neighbors	N/A (distance-based)
Decision Tree	Gini Impurity
SVM (RBF Kernel)	Hinge Loss
Naive Bayes	N/A (probabilistic / Bayes' theorem)
Random Forest	Gini Impurity
XGBoost	Binary Cross-Entropy (Log Loss)
LightGBM	Binary Cross-Entropy (Log Loss)

3. Model Baseline

- Each model is trained and tested on the same single dataset
- 5 datasets × 8 models = 40 baseline models
- Datasets: Statlog, CHD, Framingham, Heart, Stroke
- Models: Logistic Regression, KNN, Decision Tree, SVM, Naive Bayes, Random Forest, XGBoost, and LightGBM
- Evaluation metrics: Accuracy, F1 Weighted, F1 Macro, and Recall Minority
- Focus on **F1 Macro**

3. Model Baseline

F1 Macro



- **Ensemble models (RF, XGBoost, LightGBM) perform best** on **heart** and **statlog**
- **framingham**, **stroke**, and **chd** show generally **low F1 Macro** across most models
- Low F1 Macro on **stroke** and **framingham** reflects severe class imbalance

เราทำมากกว่าพีช

3. Model Baseline

Recall Minority

dataset	LogReg	KNN	DTree	SVM	NaiveBayes	RandForest	XGBoost	LightGBM
statlog	1.00	0.00	0.67	0.00	0.00	0.92	0.96	0.96
chd	1.00	0.94	0.88	1.00	1.00	1.00	1.00	1.00
framingham	1.00	1.00	0.57	0.00	1.00	0.36	0.98	1.00
heart	0.52	1.00	0.79	0.00	0.88	0.82	0.85	0.85
stroke	1.00	1.00	0.28	0.00	0.06	0.00	0.14	0.00

Low (misses patients)  High (catches patients)

- **SVM predicts negative for nearly all cases** → recall $\approx$ 0.00 across datasets
- **chd** shows **high recall but low F1 Macro** → models are biased toward the positive class
- **stroke ensemble models (RF, LightGBM) achieve recall = 0.00** — class imbalance effect

เราทำมากกว่าพีช

4. Concatenation Setup

1. Union

- Merge all features from all datasets
- Missing values
 - numeric: imputed with median
 - categorical: imputed with mode
- Scale with MinMaxScaler

2. Intersection

- Keep only shared features across datasets
- No missing values → scale with MinMaxScaler directly

Dataset Combinations

- 2-dataset combinations: $10 \text{ pairs} \times 2 \text{ strategies} = 20 \text{ configs}$
- 3-dataset combinations: $10 \text{ triplets} \times 2 \text{ strategies} = 20 \text{ configs}$
- Total: 40 concat configurations

4. Concatenation Setup

Rename Columns (Display only changed columns)

Statlog Heart Disease

- chest-pain → cp_type
- rest-bp → sbp
- serum-chol → chol
- fasting-blood-sugar → fbs
- electrocardiographic → ecg
- max-heart-rate → hr
- major-vessels → mv
- heart-disease → target

Coronary Heart Disease

- obesity → bmi
- chd → target

4. Concatenation Setup

Rename Columns (Display only changed columns)

Framingham

- Sex → sex
- education → edu
- currentSmoker → smoking_status
- cigsPerDay → num_cigs_per_day
- BPMeds → bp_status
- prevalentStroke → stroke_status
- prevalentHyp → hypertension_status
- diabetes → diabetes_status
- totChol → chol
- sysBP → sbp
- diaBP → dbp
- BMI → bmi
- heartRate → hr
- glucose → glucose_level
- TenYearCHD → target

4. Concatenation Setup

Rename Columns (Display only changed columns)

Heart Disease


- cp → cp_type
- trestbps → sbp
- restecg → ecg
- thalach → hr
- exang → angina
- ca → mv

Stroke

- id -/->
- gender → sex
- hypertension → hypertension_status
- heart_disease → heart_disease_status
- ever_married → marital_status
- Residence_type → residence_type
- avg_glucose_level → glucose_level
- stroke → target

4. Concatenation Setup

Example of 2-Dataset Concatenation



4. Concatenation Setup

Example of 2-Dataset Concatenation

Statlog-Framingham-Union Columns

age, sex, cp_type, sbp, chol, fbs, ecg, hr, angina, oldpeak, slope, mv, thal, target, edu, smoking_status, num_cigs_per_day, bp_status, stroke_status, hypertension_status, diabetes_status, dbp, bmi, glucose_level

Statlog-Framingham-Intersection Columns

age, sex, sbp, chol, hr, target

5. Concatenation Result

Dataset Combinations

	chd	framingham	heart	statlog	stroke	count
column						
target	✓		✓	✓	✓	5
age	✓		✓	✓	✓	5
sbp	✓		✓	✓		4
sex			✓	✓	✓	4
bmi	✓		✓		✓	3
chol			✓	✓		3
hr			✓	✓		3

เราทำมากกว่าพีช

5. Concatenation Result

	chd	framingham	heart	statlog	stroke	count
column						
mv			✓	✓		2
angina			✓	✓		2
cp_type			✓	✓		2
thal			✓	✓		2
ecg			✓	✓		2
oldpeak			✓	✓		2
smoking_status		✓			✓	2
fbs			✓	✓		2
glucose_level		✓			✓	2
slope			✓	✓		2
hypertension_status		✓			✓	2

เราทำมากกว่าพืช

5. Concatenation Result

	chd	framingham	heart	statlog	stroke	count
column						
stroke_status		✓				1
tobacco	✓					1
typea	✓					1
residence_type					✓	1
adiposity	✓					1
num_cigs_per_day		✓				1
marital_status					✓	1
ldl	✓					1
heart_disease_status					✓	1
famhist	✓					1

	chd	framingham	heart	statlog	stroke	count
column						
edu		✓				1
diabetes_status		✓				1
dbp		✓				1
bp_status		✓				1
alcohol	✓					1
work_type					✓	1

เราทำมากกว่าพีช

5. Concatenation Result

Number of Rows

Baseline

	dataset	rows
0	statlog	216
1	chd	369
2	framingham	3392
3	heart	241
4	stroke	4087

เราทำมากกว่าพีช

5. Concatenation Result

Number of Rows

2-Dataset Intersection

	dataset	rows
15	statlog_chd_intersection	517
16	statlog_framingham_intersection	3607
17	statlog_heart_intersection	457
18	statlog_stroke_intersection	285
19	chd_framingham_intersection	3759
20	chd_heart_intersection	545
21	chd_stroke_intersection	3990
22	framingham_heart_intersection	3632
23	framingham_stroke_intersection	7478
24	heart_stroke_intersection	296

2-Dataset Union

	dataset	rows
5	statlog_chd_union	585
6	statlog_framingham_union	3608
7	statlog_heart_union	457
8	statlog_stroke_union	4303
9	chd_framingham_union	3761
10	chd_heart_union	610
11	chd_stroke_union	4456
12	framingham_heart_union	3633
13	framingham_stroke_union	7479
14	heart_stroke_union	4328

เราทำมากกว่าพืช

5. Concatenation Result

Number of Rows

3-Dataset Intersection


	dataset	rows
35	statlog_chd_framingham_intersection	2647
36	statlog_chd_heart_intersection	698
37	statlog_chd_stroke_intersection	163
38	statlog_framingham_heart_intersection	3848
39	statlog_framingham_stroke_intersection	297
40	statlog_heart_stroke_intersection	297
41	chd_framingham_heart_intersection	2671
42	chd_framingham_stroke_intersection	7162
43	chd_heart_stroke_intersection	164
44	framingham_heart_stroke_intersection	300

3-Dataset Union

	dataset	rows
25	statlog_chd_framingham_union	3977
26	statlog_chd_heart_union	826
27	statlog_chd_stroke_union	4672
28	statlog_framingham_heart_union	3849
29	statlog_framingham_stroke_union	7695
30	statlog_heart_stroke_union	4544
31	chd_framingham_heart_union	4002
32	chd_framingham_stroke_union	7848
33	chd_heart_stroke_union	4697
34	framingham_heart_stroke_union	7720

เราทำมากกว่าพืช


5. Concatenation Result



- Out of 160 comparisons, Concat wins 73, loses 71, ties 16
- Win rate of ~46% – concat does not consistently outperform baseline
- Marginal improvement suggests simple concatenation is insufficient for generalization


5. Concatenation Result

Union vs Intersection: Generalization Comparison



- Intersection slightly outperforms Union (win rate ~49% vs ~43%)
- However, both strategies show negative mean $\Delta f1_macro$
- Imputing missing values in Union may introduce noise, hurting performance


5. Concatenation Result



- **chd** is the only dataset that benefits from concatenation (+0.06)
- **framingham** remains largely unchanged (~0.00)
- **heart**, **statlog**, and **stroke** show notable degradation – heart drops up to -0.25
- Datasets with stronger single-dataset baselines tend to suffer more from concat

เราทำมากกว่าพืช


5. Concatenation Result



- Most datasets **show negative Δrecall** – concat detects fewer disease cases than baseline
- **chd** drops most severely (-0.414 intersection, -0.352 union)
- **stroke + intersection** is the only case with improvement (+0.131)

6. Feature Importance

1. Statlog




- Baseline: mv, thal, cp_type dominant
- Concat: age becomes dominant, typea and tobacco enter — non-cardiac features override

เราทำมากกว่าพืช

6. Feature Importance


2. CHD



- Baseline: age and famhist are dominant across most models
- Concat: chol, oldpeak, thal enter the top 5 — features from other datasets interfere with CHD-specific signals

6. Feature Importance

3. Framingham




- Baseline: age and sbp consistently dominant
- Concat: chol and glucose_level replace some baseline features — moderate feature shift

เราทำมากกว่าพืช

6. Feature Importance

4. Heart




- Baseline: cp_type and mv are strong cardiac-specific signals
- Concat: hr, chol, sbp, tobacco take over — cardiac features are diluted

เราทำมากกว่าพืช

6. Feature Importance


5. Stroke



- Baseline: age dominant across all 8 models — very consistent signal
- Concat: feature importance spreads to oldpeak, ldl, sbp, cp_type — cardiovascular features from other datasets enters

เราทำมากกว่าพืช


7. Error Analysis



- **Decision Tree** degrades across all 5 datasets – most consistently harmful model
- **SVM** and **XGBoost** drop severely on chd (-0.658, -0.615) but recover on other datasets
- **KNN** and **Naive Bayes** benefit on statlog (+0.585, +0.593) – simple models adapt better
- **LogReg** improves on **heart** (+0.348) but hurts on **statlog** (-0.141)
- No single model consistently benefits from concat across all datasets

เราทำมากกว่าพืช

7. Error Analysis



- **chd** suffers the most — recall drops by -0.414 (intersection) and -0.352 (union)
- **stroke intersection** is the only case with meaningful improvement ($+0.131$)
- **statlog and framingham** show minor degradation overall
- Union consistently performs slightly better than intersection on recall minority

8. Conclusion

1. Does concatenation improve generalization?

- No clear improvement — win rate only 46% (73 wins, 71 losses out of 160 comparisons)
- Both Union and Intersection show negative mean $\Delta f1_macro$ overall

2. Which datasets benefit?

- chd — the only dataset that consistently improves with concat
- framingham — largely unchanged
- heart, statlog, stroke — degrade, with heart dropping most severely (up to -0.25)

3. Union vs Intersection?

- Intersection slightly better on win rate (~49% vs ~43%)
- Union introduces noise via imputation, hurting performance on feature-rich datasets

4. Clinical concern (Recall Minority)

- Most datasets detect fewer disease cases after concat
- chd drops most severely (-0.414)
- stroke intersection is the only meaningful improvement (+0.131)

8. Conclusion

5. Why does concat fail?

- Different feature distributions across datasets confuse the model
- Imputed missing values (Union) add noise
- Stronger single-dataset baselines are hard to beat

Limitation

- Only 5 datasets with notably different characteristics — each focuses on different clinical aspects of cardiovascular disease
- No hyperparameter tuning per concat configuration

Future Work

- Expand to more diverse and larger datasets
- Tune hyperparameters per concat configuration for fairer optimization
- Change ways to handle class imbalance

Source Code

<https://github.com/PATRECOG-RaoTamMakKwaPeudSSR/multidataset-health>